

# Overview of the ProtoDUNE Prompt Processing System (p3s)

## Design Paper and Motivation

Supporting documents and an outline of the design can be found in the FNAL DocDB 1861 (authorization required for access).

The p3s is the computing platform for protoDUNE to support Data Quality Management (DQM). Its requirements and mode of operation are different from those of a typical Workload Management System (such as used in offline production). In particular,

- there are specific ETA requirements ranging from minutes to tens of minutes for various types of processing jobs since for DQM purposes the results become stale (not actionable) rather fast
- only a portion of the data (configurable) needs to be processed; it is assumed that the DAQ and its monitoring system provide stable data taking conditions so it is not necessary to sample most of the data continuously
- if necessary at run time, in any stage of processing a portion of the data can be dropped (i.e. excluded from the DQM stream) in order to optimize throughput for critical jobs
- there is no retry/recovery mechanism for failing jobs since any substantial delay in processing a unit of data makes the result less relevant. Instead, output and error logs are recorded and used to debug jobs
- processing jobs/streams are initiated automatically and are triggered by the data arriving from DAQ
- in p3s there is no distinct data handling system for two reasons. First, it is assumed that p3s can access data either locally via a POSIX-like interface, or via XRootD interface with minimal coding and integration effort. Second, a full-fledged data handling system would introduce additional complexity, latency and potentially failure modes. In summary, data handling capabilities are kept to an absolute minimum.

## Job dispatch

### Pilots

To minimize latency and provide the ability to run transparently on a few local resources (e.g. some ad-hoc cluster at CERN, central CERN batch facility (Tier-0) and perhaps some other facilities) any reliance on the flavor of the underlying batch system needs to be eliminated. In addition, latencies inherent in any batch

system should be optimally mitigated. Both problems are addressed by utilizing a pilot-based job dispatch, where the pilots (agents) contact a central service and only receive jobs in case the batch slot is secured and the environment validated. This also combats a few failure modes.

## Pilot States

An example of what states a pilot can go through during its lifecycle is given below:

- *active*: registered on the server, no attempt at brokerage yet
- *no jobs*: no jobs matched this pilot
- *dispatched*: got a job and preparing its execution (may still fail)
- *running*: running the payload job
- *finished*: job has completed
- *stopped*: stopped after exhausting all brokerage attempts.

## Workflows

Workflow model, interface and the corresponding p3s client are described in a separate document (*WORKFLOW.pdf*).

## Location of the input raw data

The protoDUNE DAQ writes the data to its own “online buffer” from which it is transferred to CERN EOS (centralized distributed high-performance disk storage). In order to continue operation in case of a network outage which could make EOS inaccessible, the system must be able to optionally feed from the online buffer without putting too much extra I/O load on it.

## Just-in-time job execution

The near-time nature of prompt processing requires “just-in-time” job submission and not be subjected to the often unpredictable latencies found in batch systems. An efficient and tried way to achieve this is the pilot-based job submission.

## Components

- Web service:

- workflows and their templates (dags)
- jobs
- data
- handling of pilots' requests for registration and payload
- Clients
- The *pilot* - submission and management of pilot data on the server
- The *job* - submission of job definitions to the server and management of job data on the server